

L&D Shakers AI Sandbox — Storyboard v2

What this document is: *A stakeholder storyboard describing how the AI Sandbox learning experience works. It is not a course outline or a technical specification. It is a communication artifact — a way to walk someone through the experience before they are inside it.*

What changed from v1: *Three improvements have been applied throughout. (1) The opening proposition now leads with what participants get to keep, not what AI unlocks. (2) The bi-weekly session now offers two modes — Sprint and Studio — so participants choose the pace that matches their relationship with AI. (3) The Override Log has been redesigned as a live notation practice, not a retrospective form.*

The Opening Question

Most AI programmes open by telling you what AI can do.

This one opens differently.

What do you need to protect?

Your credibility in a stakeholder room. Your professional judgment. The voice in your writing that clients recognise as yours. The standards you've spent years building. The instinct that tells you when something is technically correct but fundamentally wrong.

These are not things AI can give you. They are things an AI programme — if it's designed carelessly — can quietly erode.

The L&D Shakers AI Sandbox is built on one belief: **your expertise is the thing AI cannot replace.** Everything in this programme is designed to make that expertise more visible, not less. More legible to stakeholders. More deliberate in your own practice. More present in your outputs, even when AI helped produce them.

This is not a programme about learning AI tools.

It is a programme about learning to work with AI in a way that makes you more of what you already are.

What This Is

The AI Sandbox is a **10-week, cohort-based practice space** for Learning & Development professionals who want to redesign their workflows using AI — deliberately, safely, and in a way they can defend.

Not a course. Not a certification. A studio.

In a studio, you build things. You test them. You show your process to peers, not just your finished output. You make decisions in front of people who can ask you why. You leave with artefacts you created, not slides you consumed.

The programme runs in bi-weekly cycles. Each cycle, you work on one real workflow challenge from your own practice. You choose the pace that fits your context. You document your judgment — not AI's outputs.

And at the end of 10 weeks, you have something that no AI generated: a body of work that shows exactly where your expertise made the difference.

Who This Is For

L&D professionals who have tried AI — and know there is more, but haven't yet found the right conditions to go deeper.

Specifically, two kinds of people tend to find this programme most valuable:

The professional working inside constraints. You operate in a regulated, risk-aware environment. You may have limited tool access — often Copilot only. You've seen colleagues rush into AI use without thinking through governance, and you don't want to be that person. You want to move forward, but carefully, with a clear audit trail and language you can use in a stakeholder room.

The professional who values craft. Your work is your professional identity. You've spent years developing a voice, a standard, a way of approaching learning design that clients and colleagues recognise as distinctly yours. You're not resistant to AI — you're protective of what makes your work worth having. You want AI to expand what you can do, not flatten it.

Both of these professionals belong in this programme. The design honours both orientations.

The Core Belief: Your Judgment Is the Product

Before describing the programme structure, this needs to be said plainly.

Every AI tool produces outputs. Some of them are excellent. Many are serviceable. Some are subtly wrong in ways only an experienced professional would catch.

The skill this programme develops is not prompting. It is **judgment** — knowing when to use AI, what to give it, what to keep for yourself, and when to override what it suggests.

This judgment is documented throughout the programme using the **Override Log**: a live notation practice that runs alongside everything you make. Not a form you fill in afterwards. A record you keep as you work.

The Override Log has three fields:

- **AI suggested** — what the tool produced or proposed
- **I changed** — what you modified, redirected, or replaced
- **Because** — the reasoning that only you could apply: the organisational context, the learner reality, the professional standard that AI had no way to know

Over 10 weeks, your Override Log becomes something more than a reflection tool. It becomes an **evidence base for your own expertise** — a record of every time your judgment caught something the algorithm missed. For some participants, it becomes a governance document they share with stakeholders. For others, it becomes a craft journal that shapes how they work with AI permanently.

This is the core artefact of the programme. Everything else is scaffolding around it.

The 10-Week Journey

The programme runs across 10 weeks, structured in five bi-weekly cycles followed by two culminating weeks.

```
Week 1 – Cycle 1, Session: Show & Tell + Sprint/Studio · Focus: Define
Week 2 – Cycle 1, Between: Tiny Experiment + Reflection Pod
Week 3 – Cycle 2, Session: Show & Tell + Sprint/Studio · Focus: Discover
Week 4 – Cycle 2, Between: Tiny Experiment + Reflection Pod
Week 5 – Cycle 3, Session: Show & Tell + Sprint/Studio · Focus: Design & Develop
Week 6 – Cycle 3, Between: Tiny Experiment + Reflection Pod
Week 7 – Cycle 4, Session: Show & Tell + Sprint/Studio · Focus: Deploy & Iterate
Week 8 – Cycle 4, Between: Tiny Experiment + Reflection Pod
Week 9 – Deep Dive with Peers
Week 10 – Transformation Expo
```

Each cycle moves through one focus area of the L&D workflow: Define → Discover → Design & Develop → Deploy & Iterate. By Week 8, participants have tested AI across the full arc of their professional work, not just in one corner of it.

How Each Cycle Works: The Bi-Weekly Rhythm

Each two-week cycle contains five elements. Together they take approximately 145 minutes across the fortnight — structured to fit real working lives, not ideal ones.

01 — Show & Tell · 20 minutes · Live

Each session opens with a guest practitioner sharing one real workflow transformation. Not a vendor case study. Not a best-practice overview. A specific story: here is the challenge I had, here is what I tried, here is where AI helped, here is where my judgment made the difference.

The guest's Override Log — the decisions they made, the things they changed, the reasoning behind it — is shared as part of the story. This is not incidental. It is the point.

The Show & Tell is always connected to the cycle's focus area. A session focused on Define might feature a guest who used AI to analyse 2,000 survey responses and had to override the tool's ranking of learning needs because they knew something about the organisation that no dataset could contain.

Guests are chosen for credibility and honesty, not for enthusiasm. Participants who are ambivalent about AI — who are watching carefully before committing — need to see practitioners who share that ambivalence and found a way through it.

02 — Sprint or Studio · 45 minutes · Live

This is the core active learning session of each cycle. Participants work on a real micro-case using AI. What changes in v2 is that **participants choose how they work**.

Sprint mode is for those who want to move fast and learn through doing. Blank canvas, timer running, one workflow challenge, AI tools open. The goal is a functional output produced under time pressure — and an Override Log capturing every decision made along the way. Sprint participants share outputs at the end of the session.

Studio mode is for those who need to trace a path before they walk it. The studio opens with a worked example from a previous cohort: a real experiment, with the participant's Override Log visible and narrated.

Studio participants then work on their own challenge, but with their process visible to the group — peers can see what they're making as they make it, and conversation happens mid-work, not only after. Studio participants share process as well as outputs.

Both modes use the same workflow challenge. Both produce an Override Log. Both count equally in the programme.

The choice of mode is made at the start of each session and can change cycle to cycle. Over time, participants often find they shift modes depending on the focus area — sprinting through familiar territory, slowing to studio pace when the work touches something they're protective of. That pattern itself becomes a data point about where AI feels safe and where expertise needs protecting.

03 — Tiny Experiment · 60 minutes async · Between sessions

Before leaving the live session, each participant commits to one small experiment to run in their actual workplace before the next cycle.

The goal is minimum viable, not perfect. One prompt. One workflow step. One real piece of work, tested with AI, with the Override Log running alongside it.

The Tiny Experiment is not a homework assignment. It is a commitment made in front of peers, which is a different thing. The social contract of the cohort — three people in a pod who will ask you what you shipped — is the mechanism that makes this happen. Not obligation. Accountability.

The Override Log for a Tiny Experiment is started during the work, not after. Participants are encouraged to keep a document open alongside their AI tool with three live fields: what AI suggested, what I changed, because. The quality of this notation is more important than the sophistication of the experiment. A simple, honest Override Log from a basic experiment is worth more than a complex experiment with no record of judgment.

04 — Reflection Pod · 10 minutes · Between sessions

Between sessions, each participant meets briefly with their two pod peers — synchronously or asynchronously — to share what they built and what they learned.

The three pod prompts are:

1. **What did you ship?** Show the artefact: the prompt, the output, and the Override Log.
2. **What did you change?** Walk through your override reasoning. Not just what you changed — why.
3. **What can you do now that you couldn't before?** This question is asked every cycle, not just at the end.

The pod's job is not to evaluate quality. It is to ask the questions that make judgment visible. *Why did you phrase the prompt that way? What made you override that specific word? What would have happened if you hadn't caught that?*

Pods are groups of three, assigned at the start of the programme. Three is the minimum for productive challenge and the maximum for genuine intimacy. Everything shared in a pod stays in the pod unless the participant explicitly chooses to bring it to a wider session.

05 — Override Log · Running · Throughout

The Override Log is not a separate step. It is a practice that runs through all four previous elements.

In the Show & Tell, you see someone else's Override Log and note what it surfaces for you.

In the Sprint or Studio, your Override Log is open as you work.

In the Tiny Experiment, your Override Log documents your judgment in real time.

In the Reflection Pod, your Override Log is the artefact you share.

By the end of each cycle, your Override Log for that experiment has four dimensions:

- **The live notation** made during the work
- **The pod conversation** that interrogated your reasoning
- **The pattern** you're starting to see across cycles — where you override most often, and why

This pattern is the real learning. After four cycles, participants can typically articulate — with evidence — where their expertise matters most and what kind of AI assistance genuinely improves their work versus what kind subtly degrades it.

The Experiment Archive

Each Tiny Experiment is recorded in the cohort's shared archive. The archive is not a showcase. It is a **working collection** — experiments in various states of completion, with honest Override Logs, including the ones where AI underperformed or where the participant's judgment prevented a significant error.

The archive exists for two reasons.

First, for the cohort: to learn from each other's judgment, not just each other's outputs. The most valuable entries in the archive are often the ones where the override reasoning is complicated — where the participant wasn't sure, where the stakes were real, where the decision between accepting AI's suggestion and changing it was genuinely difficult.

Second, for each participant individually: to have a record of their professional judgment developed over 10 weeks. This record is theirs to keep, use, and share as they choose. For participants working in regulated environments, it can function as a governance document demonstrating that AI use in their practice is considered, validated, and audited. For participants who value craft, it is a portfolio of professional judgment — evidence not of what AI produced, but of what they brought to the work that AI could not.

Week 9: Deep Dive with Peers

In Week 9, the cohort comes together for a focused 60-minute session to go deeper on one workflow area. Each cycle's Deep Dive rotates through the four focus areas, so across the programme participants encounter all four in depth.

The session structure is Frame → Build → Debrief.

Frame (15 min): Participants identify their specific challenge within the focus area and define what "better" looks like in their context — not in the abstract. A Define Deep Dive doesn't ask "how can AI improve needs analysis?" It asks "where in *your* needs analysis does insight arrive too late to influence decisions?"

Build (30 min): Participants work on their specific challenge with AI, Override Log open. A worked example opens the Build phase — a real experiment from a previous cohort with the full Override Log visible, including where the participant was uncertain and what they did with that uncertainty. Studio mode applies throughout: process is visible, conversation happens mid-work.

Debrief (15 min): Participants share what worked, what they overrode, and what they'd take back to their team. The debrief question is consistent across all Deep Dives: *"What did you protect in this experiment, and how will you explain that decision?"*

The Four Deep Dive Topics

Define — Needs Analysis

The guiding question: *Where does context matter more than data?*

AI can surface patterns in qualitative feedback that humans miss. It can also confidently misrank priorities because it lacks the organisational knowledge to understand why one need is urgent and another is chronic. This Deep Dive focuses on the boundary — knowing when AI's pattern recognition helps and when your contextual judgment must override it.

Discover — Research & Benchmarking

The guiding question: *What does AI miss that your expertise catches?*

AI accelerates research. It also flattens it — producing comprehensive-looking outputs that may lack the domain specificity, sectoral nuance, or source credibility that a careful professional would notice. This Deep Dive focuses on AI as a research accelerant, with human validation discipline built into the workflow rather than added afterwards.

Design & Develop — Learning Experience Design

The guiding question: *How do you keep the human touch in AI-designed experiences?*

AI can generate learning journeys, assessments, and scenarios quickly. The risk is sameness — outputs that are technically correct but experientially generic, lacking the learner-specific detail and emotional intelligence that make a learning experience land. This Deep Dive focuses on voice-preserving prompting and "AI as critic, not creator" workflows.

Deploy & Iterate — Implementation & Improvement

The guiding question: *What feedback loops matter most?*

Deploying AI-enhanced workflows requires knowing what to measure and who to trust. This Deep Dive focuses on the practical challenges of pilot and scale — including how to explain AI-assisted outputs to stakeholders, what validation protocols look credible under scrutiny, and how to design feedback loops that catch errors before they compound.

What You Leave a Deep Dive With

- A tested prompt and Override Log pair you can reuse
- A clearer sense of where your expertise is most irreplaceable in this focus area
- Peer feedback on your reasoning, not just your output
- Language to explain your AI-assisted decisions to someone who wasn't in the room

Week 10: Transformation Expo

The Transformation Expo is the programme's culminating session. It uses a World Café format — rotating conversations, each table hosted by a participant presenting one real transformation from their practice.

Format: 6 rotations × 10 minutes = 60 minutes total. Each table host presents their story once; participants rotate to new tables after each round.

What each table contains:

- A before/after account of a real workflow change

- The key prompt that made the change possible
- The Override Log entry that mattered most — the decision that required human judgment
- A conversation prompt for the group: *"Where in your workflow would this be hardest to implement, and why?"*

The Expo is not a showcase of AI capabilities. It is a **showcase of professional judgment applied to AI use**. The most memorable tables are the ones where the host can say clearly: *here is what AI got wrong, here is how I caught it, and here is what the outcome would have been if I hadn't*.

Programme Awards

The programme recognises participants who demonstrated exceptional professional judgment throughout the 10 weeks. Awards are deliberately framed around judgment, not boldness.

The Override Award — for the participant whose Override Log entry revealed the most significant gap between AI's output and what the situation actually required. This award goes to the person who caught something that mattered.

The Rigour Award — for the participant who developed the most thorough and replicable validation practice across their experiments. This award goes to the person whose process others want to borrow.

The Craft Award — for the participant who most consistently produced work where AI's contribution was invisible in the best sense — absorbed into, not substituted for, their professional voice.

The Transferability Award — for the participant whose prompt, workflow, or Override Log structure was adopted by the most cohort members. This award goes to the person who made the cohort better.

What Emerges Across All Tables

Patterns that consistently appear across cohort transformation stories:

1. **Human judgment matters most at the framing, validation, and interpretation stages** — not at the generation stage. AI produces. Humans decide.
2. **Override logging surfaces tacit knowledge** — experts often don't know what they know until they have to articulate why they changed something.
3. **The smallest experiments produce the most transferable insights** — complexity obscures the learning; simplicity reveals it.
4. **Peer interrogation of override reasoning sharpens professional judgment faster than solo reflection** — the pod conversation is where the real development happens.
5. **AI-generated outputs require domain-specific contextualisation that only the practitioner can provide** — this is not a limitation to work around. It is the value of the professional.

The Promise

At the end of 10 weeks, participants should be able to say:

"Thanks to this programme, I can now see exactly where my expertise matters — and I have the evidence to prove it."

This is different from the typical AI programme promise. It does not promise that participants will do things faster (though many will). It does not promise that AI will unlock new capabilities they never had (though

some will find this). It promises something more durable: **a practitioner who knows what to protect, what to delegate, and how to tell the difference.**

Specifically, participants leave with:

- A body of Override Logs documenting professional judgment applied across four workflow areas
 - A record of Tiny Experiments they can share selectively — with peers, with stakeholders, with teams they want to bring along
 - A prompting practice shaped around their actual work, not generic use cases
 - A peer relationship with two people who have interrogated their thinking over 10 weeks
 - Language for explaining AI-assisted decisions to stakeholders who were not in the room
 - A clearer answer to the question that matters most: *what can I do now, with AI, that I couldn't do before — and what am I doing better because I didn't let AI do it alone?*
-

Sandbox Principles

These principles are not aspirational. They are design constraints. Every element of the programme is built around them.

Your judgment is the curriculum. The programme doesn't teach you what to think about AI. It creates conditions for you to discover what you already think — and make it visible.

Imperfect thinking is the material. The most valuable moments in this programme happen when someone says *"I'm not sure I made the right call here."* That is not failure. That is the work.

Process beats output. A prompt and an Override Log together are worth more than a polished deliverable with no record of how it was made.

Constraints are context, not obstacles. If you can only use Copilot, that is your context. The programme works within your reality, not an idealised version of it.

Slow is allowed. Studio mode exists because some of the most important professional development happens at a pace that a 45-minute sprint cannot accommodate. Choose the speed that produces real learning, not the speed that produces the appearance of it.

What you protect defines what you value. The Override Log is ultimately a values document. What you consistently choose to change, and why, tells you something about the professional you are and the standards you hold. That is worth knowing.

A Note on Tool Access

The programme is designed to work within real organisational constraints. Participants are not expected to have access to a full suite of AI tools.

All programme experiments are tested against the following tool constraint: *can this be done using only Microsoft Copilot?* Where the answer is yes, the Copilot-compatible workflow is documented and available. Where more capable tools open additional possibilities, those are offered as extensions, not requirements.

Participants who work in regulated environments — where AI tool use requires approval, where outputs must be validated before use, where data privacy governs what can be entered into a prompt — are supported

with validation templates and governance framing built into the experiment structure. The Override Log itself is designed to function as a governance-ready audit trail for AI-assisted work.

How This Storyboard Is Organised

This storyboard follows the sequence a participant would experience the programme:

1. **The opening question** — what they are asked to protect
2. **The programme overview** — what kind of experience this is
3. **Who it's for** — and the honest acknowledgement that this is designed for ambivalence, not enthusiasm
4. **The core belief** — the Override Log philosophy stated before the structure
5. **The 10-week journey** — the arc at a glance
6. **The bi-weekly cycle** — how each fortnight works, including both Sprint and Studio modes
7. **The experiment archive** — what gets collected and why
8. **Week 9: Deep Dive** — going deeper with peers
9. **Week 10: Transformation Expo** — the culminating showcase
10. **The promise** — what participants leave with
11. **Sandbox principles** — the design constraints behind every decision
12. **A note on tool access** — the reality of organisational constraints

Stakeholders reviewing this storyboard are encouraged to read it in the order a participant would encounter it. The sequence matters: the opening question is designed to be heard before the structure is revealed.

Storyboard v2 — Revised to reflect three improvements: (1) programme opens with what participants protect, not what AI unlocks; (2) bi-weekly session offers Sprint and Studio modes; (3) Override Log redesigned as live notation practice throughout the programme.
